

```

import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.applications import DenseNet121
from tensorflow.keras.layers import GlobalAveragePooling2D, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical
from sklearn.model_selection import train_test_split

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_val, y_train, y_val = train_test_split(x_train, y_train, test_size=0.2, random_state=42)

    Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
    170498071/170498071 [=====] - 3s 0us/step

# Normalize pixel values to the range [0, 1]
x_train, x_val, x_test = x_train / 255.0, x_val / 255.0, x_test / 255.0

# One-hot encode labels
y_train = to_categorical(y_train, 10)
y_val = to_categorical(y_val, 10)
y_test = to_categorical(y_test, 10)

# Load pre-trained DenseNet121 model (without top layers)
base_model = DenseNet121(weights='imagenet', include_top=False, input_shape=(32, 32, 3))

    Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/densenet/densenet121\_weights\_tf\_dim\_ordering\_tf\_29084464/29084464 [=====] - 0s 0us/step

# Add custom top layers for CIFAR-10
x = GlobalAveragePooling2D()(base_model.output)
x = Dense(1024, activation='relu')(x)
predictions = Dense(10, activation='softmax')(x)

# Create the model
model = Model(inputs=base_model.input, outputs=predictions)

# Compile the model
model.compile(optimizer=Adam(learning_rate=0.0001), loss='categorical_crossentropy', metrics=['accuracy'])

# Fine-tune the model on CIFAR-10 data
history = model.fit(x_train, y_train, batch_size=64, epochs=100, validation_data=(x_val, y_val))

```



```
625/625 [=====] - 46s 73ms/step - loss: 0.0155 - accuracy: 0.9949 - val_loss: 0.8270 - val_accuracy: 0.8
Epoch 90/100
625/625 [=====] - 46s 73ms/step - loss: 0.0155 - accuracy: 0.9949 - val_loss: 0.8270 - val_accuracy: 0.8
Epoch 91/100
625/625 [=====] - 47s 75ms/step - loss: 0.0091 - accuracy: 0.9973 - val_loss: 0.8223 - val_accuracy: 0.8
Epoch 92/100
625/625 [=====] - 46s 74ms/step - loss: 0.0122 - accuracy: 0.9960 - val_loss: 0.8491 - val_accuracy: 0.8
Epoch 93/100
625/625 [=====] - 47s 75ms/step - loss: 0.0113 - accuracy: 0.9966 - val_loss: 0.9705 - val_accuracy: 0.8
Epoch 94/100
625/625 [=====] - 48s 77ms/step - loss: 0.0106 - accuracy: 0.9965 - val_loss: 0.7966 - val_accuracy: 0.8
Epoch 95/100
625/625 [=====] - 47s 76ms/step - loss: 0.0080 - accuracy: 0.9975 - val_loss: 0.8489 - val_accuracy: 0.8
Epoch 96/100
625/625 [=====] - 48s 76ms/step - loss: 0.0078 - accuracy: 0.9975 - val_loss: 0.8767 - val_accuracy: 0.8
Epoch 97/100
625/625 [=====] - 47s 75ms/step - loss: 0.0110 - accuracy: 0.9963 - val_loss: 0.8109 - val_accuracy: 0.8
Epoch 98/100
625/625 [=====] - 47s 75ms/step - loss: 0.0120 - accuracy: 0.9960 - val_loss: 0.7811 - val_accuracy: 0.8
Epoch 99/100
625/625 [=====] - 48s 77ms/step - loss: 0.0114 - accuracy: 0.9961 - val_loss: 0.8340 - val_accuracy: 0.8
Epoch 100/100
625/625 [=====] - 45s 72ms/step - loss: 0.0100 - accuracy: 0.9967 - val_loss: 0.8277 - val_accuracy: 0.8
```

```
# Evaluate the model on the test data
test_loss, test_accuracy = model.evaluate(x_test, y_test)
print(f'Error Rate: {(1-test_accuracy) * 100:.2f}%')
```

```
313/313 [=====] - 5s 14ms/step - loss: 0.8278 - accuracy: 0.8574
Error Rate: 14.26%
```

```
# You can also save the model for future use
model.save('densenet_cifar10.h5')
```

```
/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3000: UserWarning: You are saving your model as an HDF5 file via
saving_api.save_model(
```